

AWS SAM Template Walkthrough for JEDx CAR Service

Top-level metadata

- `AWSTemplateFormatVersion / Transform (AWS::Serverless-2016-10-31)` — declares this as a SAM template so you can use shorthand (`AWS::Serverless::*`) and SAM CLI.
- `Description` — human-readable summary.
- `Globals.Function` — default settings applied to all Lambda functions unless overridden (here: `Timeout: 3s`, structured JSON logs).

Parameters

- `Prefix` — string used to prefix every resource name (e.g., `jedx-car-bucket`).
- `CollectorUrl` — URL for an external “collector” API (currently not used).

App grouping & monitoring

- `ApplicationResourceGroup (AWS::ResourceGroups::Group)` — groups all stack resources for easier discovery/monitoring.
- `ApplicationInsightsMonitoring (AWS::ApplicationInsights::Application)` — turns on CloudWatch Application Insights for the above resource group (auto-config enabled).

Storage & streams (CAR service)

- `JedxCarInputBucket (S3)` — landing bucket for incoming files (versioned). Triggers `S3ToKinesisFunction` on object creation.
- `JedxCarBucket (S3)` — main application bucket for validated/processed artifacts (versioned).
- `JedxCarErrorBucket (S3)` — captures validation failures or error payloads (versioned).
- `JedxCarInputStream (Kinesis)` — stream (`ShardCount: 1`) carrying records from S3 into the validation pipeline.

Public API (Gateway + Cognito auth)

- `JedxCarApi (AWS::Serverless::Api)`
- `Name/Stage: ${Prefix}-car-api, stage Prod`.
- `CORS`: allows `DELETE,GET,HEAD,PUT,POST` from any origin with common headers.
- `Auth`: Cognito User Pool authorizer (`JedxUserPoolAuthorizer`); Authorization header expected; reauth disabled.

Lambda: API handlers

- `JedxCarApiFunction (car_api/app.lambda_handler)`
- `Event`: ANY `/car` on `JedxCarApi` (generic endpoint).
- `Env`: `KINESIS_STREAM_NAME, DDB_TABLE_NAME`.
- `Policies`: `kinesis:PutRecord` to input stream; basic `DynamoDB` write/query on the table.
- `JedxCarRecordApiFunction (car_api/record.lambda_handler)`
- `Event`: ANY `/car/record/{senderId}/{object_type}/{RefId}`.
- `Env`: `S3_CAR_BUCKET, DDB_TABLE_NAME`.
- `Policies`: `S3` read/write on main bucket; `DynamoDB` put/update/get/query.
- `JedxCarRecordsApiFunction` (same handler as above)
- `Event`: ANY `/car/records/{senderId}` (list multiple records).
- `Env/Policies`: similar `S3 + DynamoDB` access.
- `JedxCarCollectorSendApiFunction (car_api/collector.lambda_handler)`
- `Event`: ANY `/car/collector/send` (forwards to external collector).
- `Env`: `S3_CAR_BUCKET, DDB_TABLE_NAME, COLLECTOR_URL`.
- `Note`: `COLLECTOR_URL` is hard-coded to a specific API URL; you’ve commented out the `!Ref CollectorUrl`. If you want per-env flexibility, switch to the parameter (uncomment the `!Ref` and remove the hardcoded lines).

- Policies: S3 read/write on main bucket; DynamoDB put/update/get/query.
- JedxCarUserApiFunction (car_api/user.lambda_handler)
- Event: POST /car/login (auth/user ops).
- Env: DDB_TABLE_NAME.
- Policies: DynamoDB put/update/get/query.

Lambda: data pipeline

- S3ToKinesisFunction (s3_to_kinesis_lambda/app.lambda_handler)
- Trigger: s3:ObjectCreated:* on JedxCarInputBucket.
- Purpose: reads new S3 objects and pushes a record into JedxCarInputStream.
- Policies: s3:GetObject on input bucket; kineses:PutRecord on stream.
- CarValidationFunction (car_validation_lambda/app.lambda_handler)
- Trigger: Kinesis (stream = JedxCarInputStream, BatchSize: 10, StartingPosition: LATEST).
- Purpose: validates incoming records and writes results.
- Env: S3_CAR_BUCKET, S3_CAR_ERROR_BUCKET, DDB_TABLE_NAME.
- EventInvokeConfig: MaximumRetryAttempts: 0 (no async retries—bad records won't be retried).
- Policies: s3:PutObject to the main and error buckets; DynamoDB PutItem to persist validation outcomes.

Database

- JedxCarTable (DynamoDB)
- Name: \${Prefix}-car-table.
- Keys: pk (HASH) + sk (RANGE) — typical single-table design for flexible access patterns.
- Billing: PAY_PER_REQUEST (on-demand).

Outputs

- JedxCollectorFunction — ARN of the primary API Lambda (for quick reference/integration).
- JedxCollectorFunctionIamRole — implicit role ARN created for that function (handy for attaching extra perms).

How the pieces flow

- 1) Files in → You upload to JedxCarInputBucket → S3ToKinesisFunction fires → sends a record to JedxCarInputStream.
- 2) Stream → validate → store → CarValidationFunction consumes Kinesis, writes good results to JedxCarBucket (and/or DynamoDB), errors to JedxCarErrorBucket.
- 3) API surface → JedxCarApi exposes /car, /car/record/..., /car/records/{senderId}, /car/collector/send, /car/login secured by Cognito; each path maps to the corresponding Lambda handler above.

Suggestions

- Parameterize COLLECTOR_URL: replace the hardcoded value with !Ref CollectorUrl for per-environment deploys.
- Least-privilege S3 policies: if possible, scope S3 ARNs down to paths/prefixes actually used.
- Retries on validation: consider a small MaximumRetryAttempts (e.g., 1–2) or a DLQ to avoid silent drops.